

Artificial Intelligence Homework 3

Reinforcement

2014 / 12 / 10

Question 1 – Value Iteration Agent

- An MDP is given
- $U(s)$: `self.values = util.Counter()` – a dictionary
- `__init__(self, mdp, discount = 0.9, iterations = 100)`:
For every iteration, for every state in the MDP, find the maximum value of $Q(s, a)$ for all possible actions of state s
 - Recall that $U(s) = \max_{a \in A(s)} Q(s, a)$
- `getValue(self, state)`:
return `self.values[state]`

Question 1 – Value Iteration Agent

- `getQValue(self, state, action):`

Use `getTransitionStatesAndProbs` in `mdp.py`

$$Q(s, a) = \sum_{s'} P(s' | s, a) \cdot [R(s' | s, a) + \gamma U(s')]$$

- `getPolicy(self, state):`

If terminal state, return `none`.

Else, return the action that results in the maximum value of

$$E[\text{utility of taking } a] = \sum_{s'} P(s' | s, a) \cdot U(s')$$

Question 2 – Value Iteration Agent

- Change only one of the parameters, the discount factor γ or the noise level, so that the agent will cross the bridge in the optimal policy
 - Noise level: the uncertainty of taking an action


Ex: When noise=0, for any given state s and action a in $A(s)$, there will be one s' such that $P(s'|s,a)=1$; for any other state $s'' \neq s'$, it holds that $P(s''|s,a)=0$.
 - Discount factor: the level of importance of the future rewards
- The result should be something like this:



Question 3 – Value Iteration Agent

- Adjust the parameters, including the discount factor γ , the noise level, and the living reward, so that the agent acts as the descriptions
 - Living reward: The amount of reward given when the agent is still alive (i.e. doesn't fall over the cliff)

Question 3 – Value Iteration Agent



Question 4~7 – Q Learning Agent

- Motivation: the transition probability and the reward of any given state are not known in advance.
- Construct a two dimensional (for states and actions) table to learn the utility of all states and the optimal policy.
 - One viable way to do this is to construct a “dictionary of dictionary” in python.

Question 4~7 – Q Learning Agent

- `__init__(self, **args):`
Construct your Q table here.
- `getQValue(self, state, action):`
If the state is already seen, return `Qtable(state, action)`
Otherwise, construct a new 1D array in the Qtable and set all the elements to 0
i.e.
new `Qtable(state, *)`
for all *action* in $A(state)$ set `Qtable(state, action)=0`

Question 4~7 – Q Learning Agent

- `getValue(self, state)`:
If there are no legal actions, return 0
Otherwise, return $\max_{action \text{ belongs to } A(state)} Qtable(state, action)$
 - Note: Please be advised to use the function “getQValue” instead of directly accessing the data in the Qtable here.
- `getPolicy(self, state)`:
- `getAction(self, state)`:
- `update(self, state, action, nextState, reward)`:
Too simple to allow any hints...

Question 9 – Approximate Q Learning Agent

- Motivation: the original Q learning method is not scalable.
- Extract the features of the state-action pair and learn the “weights” of the features instead.
- You only have to initialize the weights (you can use `util.Counter`) and override two functions “`getQValue`” and “`update`” according to the equations in the html file. You might need to call the function “`getFeatures`” defined in “`featureExtractors.py`”.